

Core AI Engineering

A Complete Guide

10 chapters compiled into one PDF

Snehal Kadiya

+91 9974031480

snehaliteng@gmail.com

<https://snehaliteng.github.io/>

Copyright & Disclaimer

© 2026 Snehal Kadiya. All rights reserved.

This tutorial is intended for personal learning and educational purposes only. No part of this document may be reproduced, distributed, or transmitted in any form or by any means without the prior written permission of the author.

Please do not print this document unnecessarily. Think of the environment — save paper and trees. Share the digital link instead:
<https://snehaliteng.github.io/tutorials/ai-engineer-core/>

[← Back to Tutorials](#)

Build Your First LLM Product

Running LLMs Locally with Ollama

Ollama lets you run open-source LLMs locally. Install it, pull a model, and make your first API call.

```
# Install Ollama (macOS/Linux)
curl -fsSL https://ollama.com/install.sh | sh

# Pull a model
ollama pull llama3.2


# Python API call
import ollama
response = ollama.chat(model="llama3.2", messages=[
    {"role": "user", "content": "Summarize a website about AI."}
])
print(response["message"]["content"])
```

LLM Engineering Building Blocks

Block	Description
Models	Base, Chat, and Reasoning variants
Tools	Function calling, code execution, search
Techniques	Prompt engineering, RAG, fine-tuning

Understanding LLM Types

- **Base Models** — Predict next token, need instruction tuning
- **Chat Models** — Fine-tuned for conversation
- **Reasoning Models** — Chain-of-thought, math, logic (e.g., o1, DeepSeek-R1)



Transformers Architecture

From LSTMs to Transformers — the key innovation is self-attention, which lets the model weigh the importance of all tokens in the input simultaneously.

```
# Key parameters
# - Parameters: 7B, 13B, 70B (more = more capability)
# - Context Window: 4K, 8K, 128K tokens
# - Tokenization: subword encoding (tiktoken, SentencePiece)
```

Business Applications

Build a sales brochure generator, chain multiple LLM calls, and stream results for real-time UX.

```
# Streaming example
for chunk in ollama.chat(model="llama3.2", messages=[...], stream=True):
    print(chunk["message"]["content"], end="", flush=True)
```

 **Exercise:** Install Ollama, pull the Llama 3.2 model, and build a website summarizer. Given a URL, fetch the page content and use the LLM to generate a 3-sentence summary. Stream the response token by token.

[← Back to Tutorials](#)

Build a Multi-Modal Chatbot

Multi-Model Conversations

Use LangChain or LiteLLM to switch between frontier models (GPT-4o, Claude, Gemini) with the same interface.

```
pip install langchain langchain-openai litellm
```

Gradio UIs

Gradio lets you prototype AI interfaces in minutes.

```
import gradio as gr

def chat(message, history):
    return f"You said: {message}"

gr.ChatInterface(chat, title="AI Chatbot").launch()
```

System Prompts & Multi-Shot Prompting

```
system_prompt = "You are a helpful airline assistant."
messages = [
    {"role": "system", "content": system_prompt},
    {"role": "user", "content": "Book a flight from NYC to London."}
]
```

Tool Calling with Airline Assistant

Build an agentic workflow where the LLM calls tools to query an SQLite database for flights.

```
import sqlite3

def search_flights(origin, destination, date):
    conn = sqlite3.connect("flights.db")
    cursor = conn.cursor()
    cursor.execute("SELECT * FROM flights WHERE origin=? AND destination=?",
                  (origin, destination))
    return cursor.fetchall()
```

Multi-Modal Capabilities

- **DALL-E 3** — Generate images from text descriptions
- **Text-to-Speech** — Convert responses to spoken audio
- **Gradio Blocks** — Build complex UIs with tabs, rows, and columns

 **Exercise:** Build a multi-modal airline assistant chatbot with: (1) a Gradio ChatInterface, (2) SQLite tool for flight searches, (3) DALL-E integration for destination images, and (4) TTS for spoken responses. Use system prompts to guide the conversation.

[← Back to Tutorials](#)

Open-Source Gen AI with HuggingFace

The HuggingFace Platform

HuggingFace hosts 500K+ models, 100K+ datasets, and Spaces for demo deployment. Use Colab with free GPUs for inference.

```
pip install transformers torch diffusers accelerate
```

HuggingFace Pipelines

Quick inference for common tasks without boilerplate.

```

from transformers import pipeline

# Sentiment
classifier = pipeline("sentiment-analysis")
result = classifier("I love AI agents!")
print(result)

# NER
ner = pipeline("ner")
print(ner("John works at Microsoft in Redmond."))

# Question Answering
qa = pipeline("question-answering")
result = qa(question="What is AI?", context="AI is the simulation of human in")
print(result)

# Image generation
from diffusers import StableDiffusionPipeline
pipe = StableDiffusionPipeline.from_pretrained("runwayml/stable-diffusion-v1-5")
image = pipe("A robot writing code").images[0]

```

Tokenizers

Tokenizers convert text to token IDs and back. Different models use different tokenizers.

```

from transformers import AutoTokenizer

tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")
tokens = tokenizer("Hello, how are you?", return_tensors="pt")
print(tokenizer.decode(tokens["input_ids"][0]))

```

Transformers & Quantization

Use 8-bit or 4-bit quantization to run large models on limited hardware.

```
from transformers import BitsAndBytesConfig

quant_config = BitsAndBytesConfig(load_in_4bit=True)
model = AutoModelForCausalLM.from_pretrained(
    "meta-llama/Llama-2-7b-chat-hf",
    quantization_config=quant_config
)
```

Applications

- **Token-by-token visualization** — See how models generate text
- **Meeting Minutes with Whisper** — Transcribe audio to text
- **Synthetic Data Generation** — Use LLMs to create training datasets

 **Exercise:** Use HuggingFace pipelines to build a meeting assistant: (1) transcribe an audio file with Whisper, (2) run NER to extract names and dates, (3) generate a summary with a text generation model, and (4) visualize the token-by-token generation process.

LLM Showdown

Model Selection Strategy

Scaling laws show that larger models (more parameters + more data) perform better, but diminishing returns apply. Use benchmarks and leaderboards to compare.

Model Size	Use Case	Cost
1-3B	Edge devices, classification	Free (local)
7-13B	Chat, summarization, RAG	Low
70B+	Complex reasoning, code gen	Medium-High
Frontier (GPT-4o, Claude)	Enterprise, multimodal	High

Commercial Use Cases

- **Automation** — Replace repetitive human tasks
- **Augmentation** — Assist humans with analysis, writing, coding
- **Agentic AI** — Autonomous decision-making systems

Code Generation Showdown

```
# Test: Convert Python to C++
prompt = "Convert this Python function to C++: def add(a, b): return a + b"

# Compare outputs from GPT-4o, Claude, DeepSeek Coder, Qwen
models = {
    "gpt-4o": "OpenAI",
    "claude-3-5-sonnet": "Anthropic",
    "deepseek-coder": "DeepSeek",
    "qwen2.5-coder": "Qwen (local via Ollama)"
}
```

Evaluation Metrics vs Outcomes


Traditional metrics (BLEU, ROUGE) don't capture real-world quality. Use LLM-as-Judge, human evaluation, and task-specific tests.

 **Exercise:** Create a model comparison test: take 5 Python functions and ask 3 different models (GPT-4o, Claude, and a local model via Ollama) to translate them to Rust. Compare the outputs for correctness, efficiency, and readability. Build a Gradio UI to display results side by side.

Mastering RAG

Fundamentals

RAG grounds LLM responses in retrieved data, reducing hallucinations and improving accuracy. Key components: embedding model, vector store, and LLM.



LangChain & Vector Databases

```
pip install langchain chromadb sentence-transformers
```

Chunking and Indexing

```
from langchain.text_splitter import RecursiveCharacterTextSplitter
from langchain_community.vectorstores import Chroma
from langchain_community.embeddings import HuggingFaceEmbeddings

splitter = RecursiveCharacterTextSplitter(chunk_size=500, chunk_overlap=50)
chunks = splitter.split_text(document)

embeddings = HuggingFaceEmbeddings(model_name="all-MiniLM-L6-v2")
vectorstore = Chroma.from_texts(chunks, embeddings)
```

Complete RAG Pipeline

```
from langchain.chains import RetrievalQA

qa = RetrievalQA.from_chain_type(
    llm=llm,
    retriever=vectorstore.as_retriever(search_kwargs={"k": 3})
)
result = qa.invoke("What is the capital of France?")
print(result)
```

RAG Evaluations

- **Faithfulness** — Does the answer match the retrieved context?
- **Relevance** — Is the retrieved context relevant to the query?
- **LLM-as-Judge** — Use an LLM to rate answer quality

Advanced RAG Techniques

- **Pre-processing** — Clean documents before chunking
- **Re-ranking** — Re-rank retrieved chunks with a cross-encoder
- **Query Expansion** — Generate multiple queries from one question
- **GraphRAG** — Build a knowledge graph from documents for better retrieval

 **Exercise:** Build a complete RAG pipeline with conversation history. Use LangChain, ChromaDB, and a Gradio chat interface. Implement: (1) document upload and chunking with overlap, (2) vector embedding and indexing, (3) hybrid search (keyword + vector), (4) conversation-aware retrieval, and (5) an LLM-as-Judge evaluation step.

[← Back to Tutorials](#)

ML to DL to Fine-Tuning

Training & Datasets

Start with a capstone project using Amazon product data from HuggingFace datasets.

```
from datasets import load_dataset

dataset = load_dataset("amazon_polarity", split="train[:1000]")
print(dataset[0])
```

Baseline Models with Traditional ML

```
from sklearn.feature_extraction.text import TfidfVectorizer
from xgboost import XGBClassifier

vectorizer = TfidfVectorizer(max_features=5000)
X_train = vectorizer.fit_transform(train_texts)
model = XGBClassifier()
model.fit(X_train, train_labels)
```

Neural Networks with PyTorch

```
import torch
import torch.nn as nn

class SentimentClassifier(nn.Module):
    def __init__(self, vocab_size, embed_dim):
        super().__init__()
        self.embedding = nn.Embedding(vocab_size, embed_dim)
        self.fc = nn.Linear(embed_dim, 2)

    def forward(self, x):
        x = self.embedding(x).mean(dim=1)
        return self.fc(x)
```

Fine-Tuning Frontier Models

Supervised Fine-Tuning (SFT) adapts a pre-trained model to your specific task. Use LoRA for efficient fine-tuning.

```
from transformers import AutoModelForCausalLM, TrainingArguments
from trl import SFTTrainer

model = AutoModelForCausalLM.from_pretrained("microsoft/Phi-3-mini-4k-instruct")
trainer = SFTTrainer(
    model=model,
    train_dataset=dataset,
    args=TrainingArguments(output_dir="./phi3-finetuned", per_device_train_batch_size=1)
)
trainer.train()
```

Monitoring & Handling Failures

- Track loss curves to detect overfitting
- Use gradient checkpointing for memory efficiency
- Save checkpoints every N steps
- Resume from checkpoints on failure

 **Exercise:** Fine-tune a small language model (Phi-3 mini or Llama 3.2 1B) on a custom dataset of your choice. Start with an XGBoost baseline, then build a PyTorch classifier, then SFT a pre-trained model. Compare all three approaches on the same test set.

Fine-Tuned Open-Source Models

QLoRA Basics

QLoRA (Quantized Low-Rank Adaptation) makes fine-tuning large models accessible by combining 4-bit quantization with low-rank adapters.

```
pip install bitsandbytes peft trl datasets
```

Loading a Quantized Model

```
from transformers import AutoModelForCausalLM, BitsAndBytesConfig
from peft import LoraConfig, get_peft_model

quant_config = BitsAndBytesConfig(
    load_in_4bit=True,
    bnb_4bit_quant_type="nf4",
    bnb_4bit_compute_dtype=torch.float16
)

model = AutoModelForCausalLM.from_pretrained(
    "meta-llama/Llama-2-7b-chat-hf",
    quantization_config=quant_config
)

lora_config = LoraConfig(
    r=16,
    lora_alpha=32,
    target_modules=["q_proj", "v_proj"],
    lora_dropout=0.05
)

model = get_peft_model(model, lora_config)
```

Dataset Preparation

Optimize token length, format as chat templates, and split train/validation.

```
def format_example(example):
    return {
        "text": f"<|user|>\n{example['question']}\n<|assistant|>\n{example['answer']}"
    }

dataset = dataset.map(format_example)
```

Hyperparameter Configuration

Use Weights & Biases (W&B) for experiment tracking and HuggingFace TRL for the trainer.

```
from trl import SFTTrainer

trainer = SFTTrainer(
    model=model,
    train_dataset=dataset,
    args=TrainingArguments(
        output_dir="./lora-llama",
        per_device_train_batch_size=2,
        gradient_accumulation_steps=4,
        learning_rate=2e-4,
        logging_steps=10,
        save_steps=100,
        report_to="wandb"  # Track with W&B
    )
)
```

Monitoring Training

Watch for loss convergence, learning rate scheduling, and overfitting via validation loss.

Inference & Evaluation

Compare fine-tuned model outputs against the base model and frontier models on test prompts.

 **Exercise:** Fine-tune Llama 3.2 1B or Phi-3 mini using QLoRA on a domain-specific dataset (e.g., medical Q&A, legal documents, code generation). Track training with W&B. Evaluate the fine-tuned model against the base model and a frontier model (GPT-4o) on 10 test prompts. Report metrics and qualitative comparisons.

[← Back to Tutorials](#)

Autonomous Multi-Agent Systems

Agentic AI & Deployment

Deploy agents to serverless cloud platforms like Modal for scalable, pay-per-use inference.

```
pip install modal
```

```
import modal

app = modal.App("agent-app")

@app.function(gpu="A100")
def run_agent(prompt: str):
    # Load model and generate
    return response


@app.local_entrypoint()
def main():
    print(run_agent.remote("Analyze this data"))
```

Advanced RAG with ChromaDB

Use ensemble retrieval with multiple embedding models for better coverage.

```
from langchain.retrievers import EnsembleRetriever

retriever = EnsembleRetriever(
    retrievers=[chroma_retriever, bm25_retriever],
    weights=[0.7, 0.3]
)
```



Structured Outputs with Pydantic

```
from pydantic import BaseModel
from typing import List

class Deal(BaseModel):
    title: str
    price: float
    discount: float
    url: str
    category: str

class DealScanResult(BaseModel):
    deals: List[Deal]
    total_savings: float
```

Planning Agents

Build agents that plan, execute, and evaluate — creating a Plan → Execute → Review loop.

```
class PlanningAgent:
    def plan(self, task):
        return self.llm(f"Create a step-by-step plan for: {task}")

    def execute(self, plan):
        for step in plan:
            result = self.execute_step(step)
            yield result

    def review(self, results):
        return self.llm(f"Evaluate these results: {results}")
```

Final Project: DealAgentFramework

Build a price-is-right deal scanner agent with:

- Web scraping MCP server for deal collection
- RAG agent for deal analysis and comparison
- Planning agent for deal monitoring schedule
- Gradio dashboard for visualization
- Modal deployment for 24/7 operation

 **Exercise:** Build a DealAgentFramework: (1) create an MCP server that scrapes deal websites, (2) a RAG agent that stores and retrieves deals, (3) a planning agent that monitors prices and alerts on drops, (4) use Pydantic for structured deal output, and (5) deploy to Modal with a Gradio frontend.

[← Back to Tutorials](#)

9. Practice & Resources

Quick Reference Guide

Task	Library / Tool	Command
Run local LLM	Ollama	<code>ollama pull llama3.2</code>
Chat completions	openai / litellm	<code>litellm --model gpt-4o</code>
Vector search	ChromaDB	<code>Chroma.from_texts(chunks, embeddings)</code>
RAG pipeline	LangChain	<code>RetrievalQA.from_chain_type(...)</code>
Fine-tuning	TRL + PEFT	<code>SFTTrainer(model, dataset)</code>
QLoRA	bitsandbytes	<code>load_in_4bit=True</code>
UI prototyping	Gradio	<code>gr.ChatInterface(...).launch()</code>
Serverless deploy	Modal	<code>modal deploy app.py</code>

Useful Resources

- [Ollama](#) — Local LLM runner
- [HuggingFace](#) — Models, datasets, Spaces
- [LangChain](#) — RAG and agent framework
- [Gradio](#) — ML demo interfaces
- [Modal](#) — Serverless GPU compute
- [TRL \(Transformer Reinforcement Learning\)](#)

8-Week Study Plan

Week	Topic	Activities
1	LLM Product	Ollama, LLM types, streaming, brochure generator
2	Multi-Modal Chatbot	Gradio, tool calling, DALL-E, TTS, SQLite
3	HuggingFace	Pipelines, tokenizers, quantization, Whisper
4	LLM Showdown	Model comparison, benchmarks, code gen
5	RAG	LangChain, ChromaDB, evaluations, advanced RAG
6	ML/DL/Fine-Tuning	XGBoost, PyTorch, SFT, monitoring
7	Open-Source Fine-Tuning	QLoRA, adapters, W&B, evaluation
8	Multi-Agent Systems	Planning agents, Modal, deal scanner capstone

Certificate Path

- **AI-900:** Azure AI Fundamentals
- **AI-102:** Azure AI Engineer Associate
- **HuggingFace Course** — Free NLP and diffusion model training

 **Final Capstone:** Combine everything into a production-ready system: (1) fine-tune a small model on your domain data using QLoRA, (2) deploy it with Modal, (3) build a RAG pipeline around it with ChromaDB, (4) add a Gradio chat interface, (5) implement multi-agent orchestration with planning and review, and (6) monitor with W&B logging.

10. References

Glossary

Term	Definition
LLM	Large Language Model — neural network trained on vast text data
Transformer	Neural architecture using self-attention
Tokenization	Converting text to subword token IDs
Embedding	Dense vector representation of text
RAG	Retrieval-Augmented Generation — ground output in retrieved data
SFT	Supervised Fine-Tuning — training on labeled examples
QLoRA	Quantized Low-Rank Adaptation — efficient fine-tuning
Quantization	Reducing model precision (32-bit to 4-bit)
Agent	AI system that uses tools, memory, and planning autonomously

Cheatsheet

```
# Installation Cheatsheet

pip install ollama                # Local LLMs
pip install langchain chromadb    # RAG
pip install transformers torch    # HuggingFace
pip install bitsandbytes peft trl # Fine-tuning
pip install gradio                # UI
pip install modal                 # Serverless

# Quick inference (HuggingFace)
from transformers import pipeline
pipe = pipeline("text-generation", model="gpt2")
pipe("Once upon a time")

# Quick RAG (LangChain + Chroma)
from langchain.chains import RetrievalQA
qa = RetrievalQA.from_chain_type(llm=llm, retriever=vectorstore.as_retriever())

# Quick fine-tuning (TRL)
from trl import SFTTrainer
trainer = SFTTrainer(model=model, train_dataset=dataset)
trainer.train()
```

Model Comparison

Model	Size	Best For	Access
GPT-4o	Frontier	General, multimodal, coding	API
Claude 3.5	Frontier	Analysis, writing, safety	API
Llama 3.2	1B-90B	Open-source, local deploy	Ollama / HF
Phi-3 / Phi-4	3.8B-14B	Efficient, edge devices	HF / Azure
DeepSeek V3 / R1	67B-671B	Reasoning, code, math	API / Ollama
Qwen 2.5	0.5B-72B	Multilingual, code	Ollama / HF

 **Congratulations!** You've completed the AI Engineer Core curriculum. You can now build LLM products, multi-modal chatbots, RAG pipelines, fine-tune models, and deploy multi-agent systems to production.